



TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Apache Spark overview


a unified analytics engine for large-scale data processing

Viet-Trung Tran

School of Information and Communication Technology

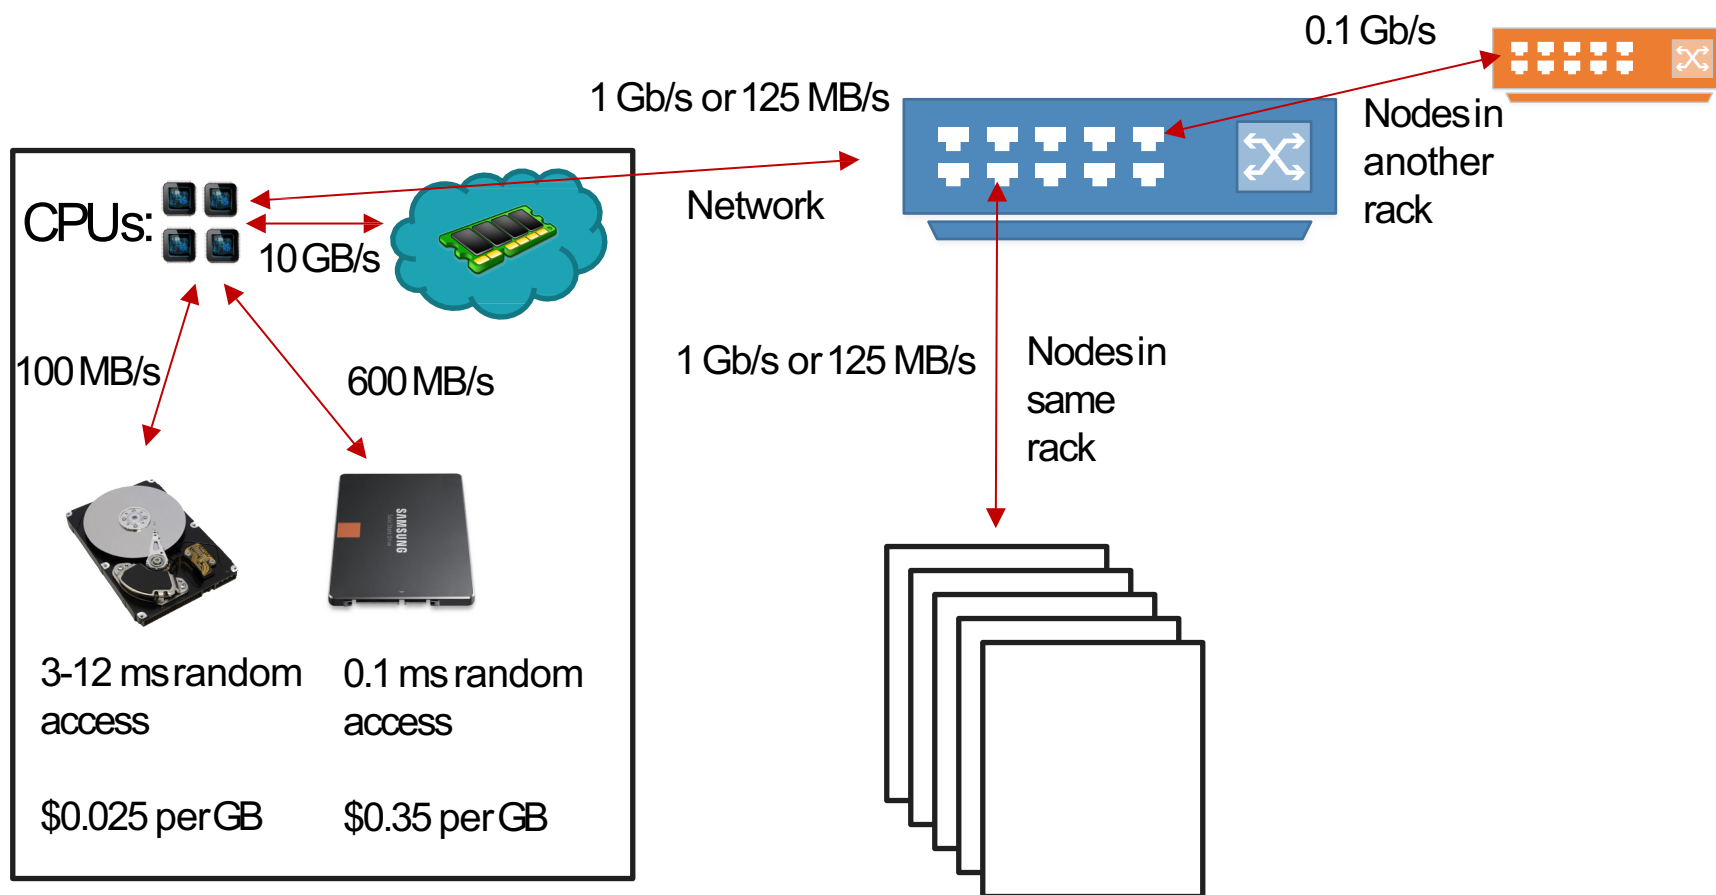
Map Reduce: Iterative jobs

- Iterative jobs involve a lot of disk I/O for each repetition




- ➔ Disk I/O is very slow!

I/O landscape

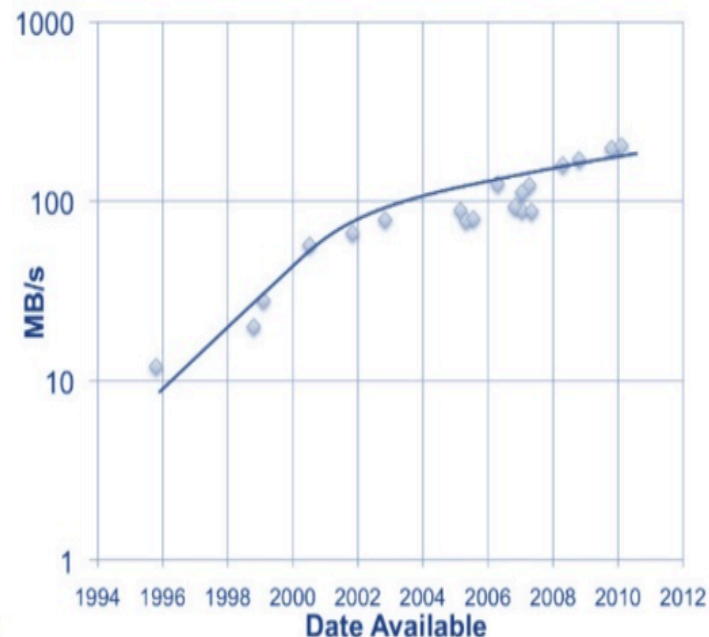


RAM is the new disk

- RAM throughput increasing **exponentially**



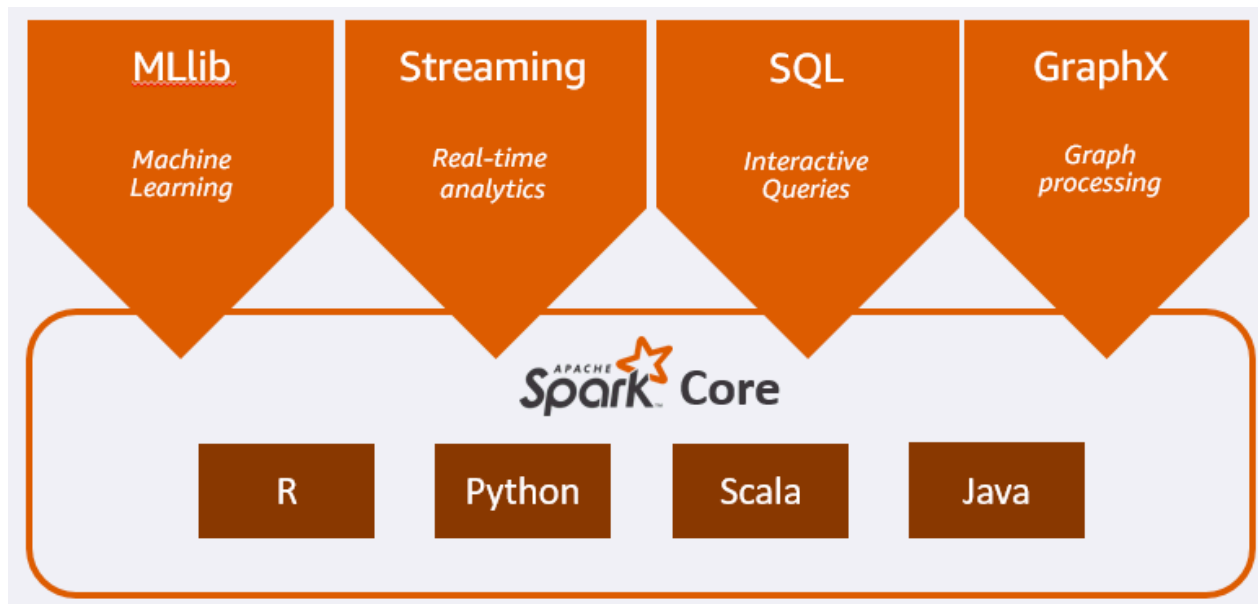
- Disk throughput increasing **slowly**



Memory-locality key to interactive response times

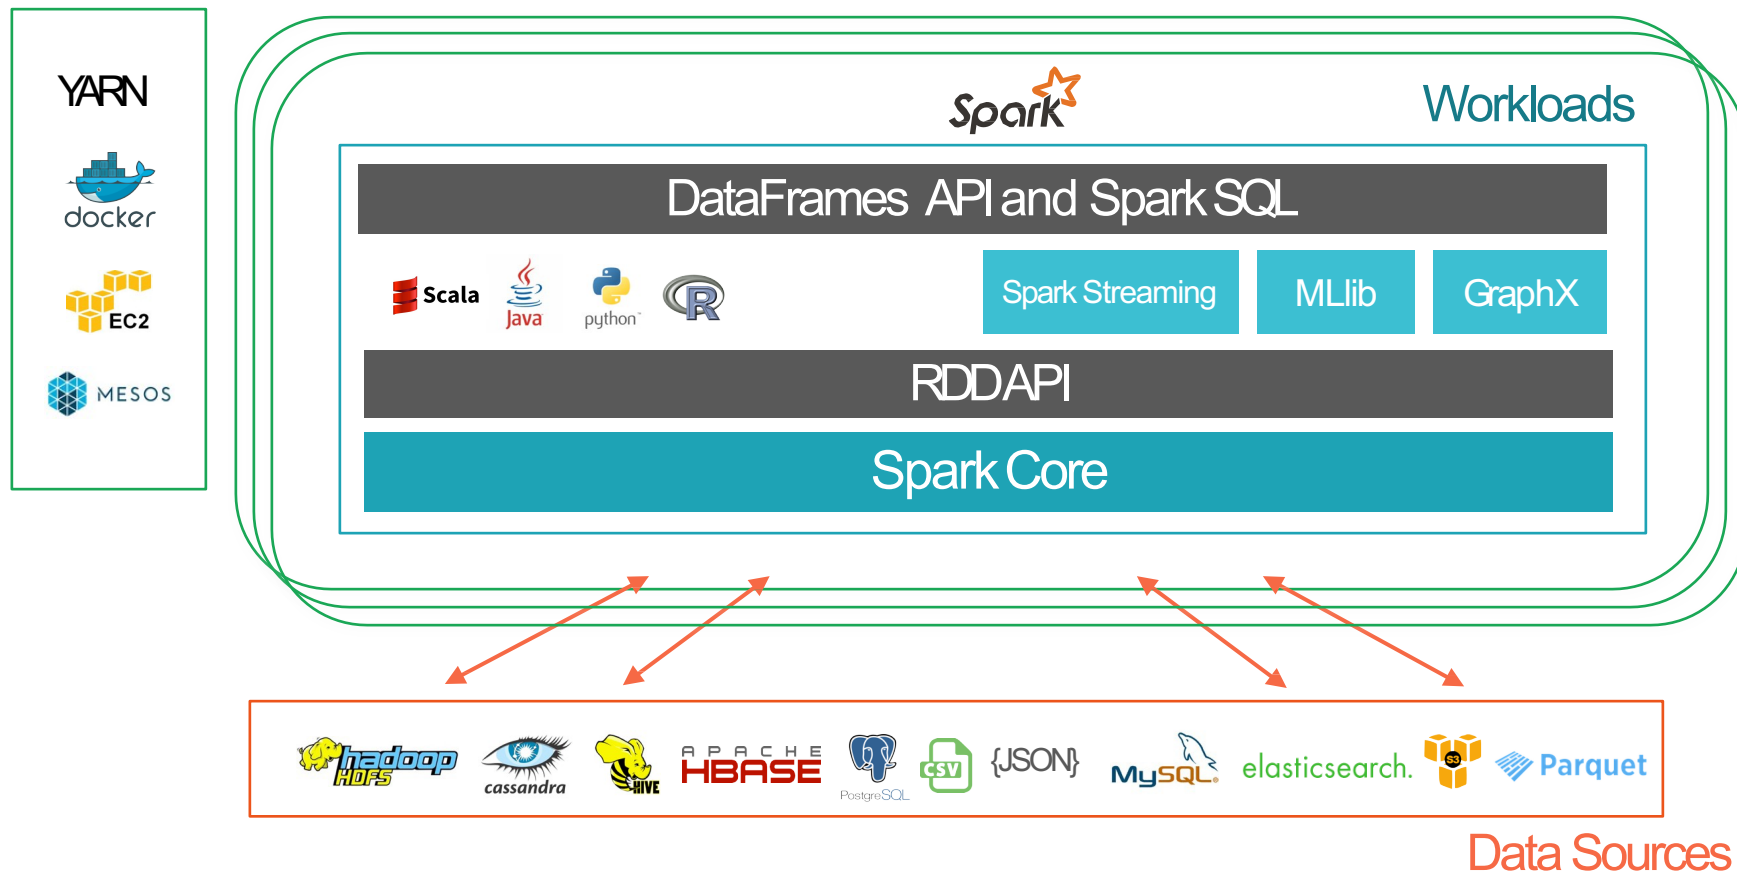
A unified analytics engine for large-scale data processing

- Better support for
 - Iterative algorithms
 - Interactive data mining
- Fault tolerance, data locality, scalability
- Hide complexities: help users avoid the coding for structure the distributed mechanism.

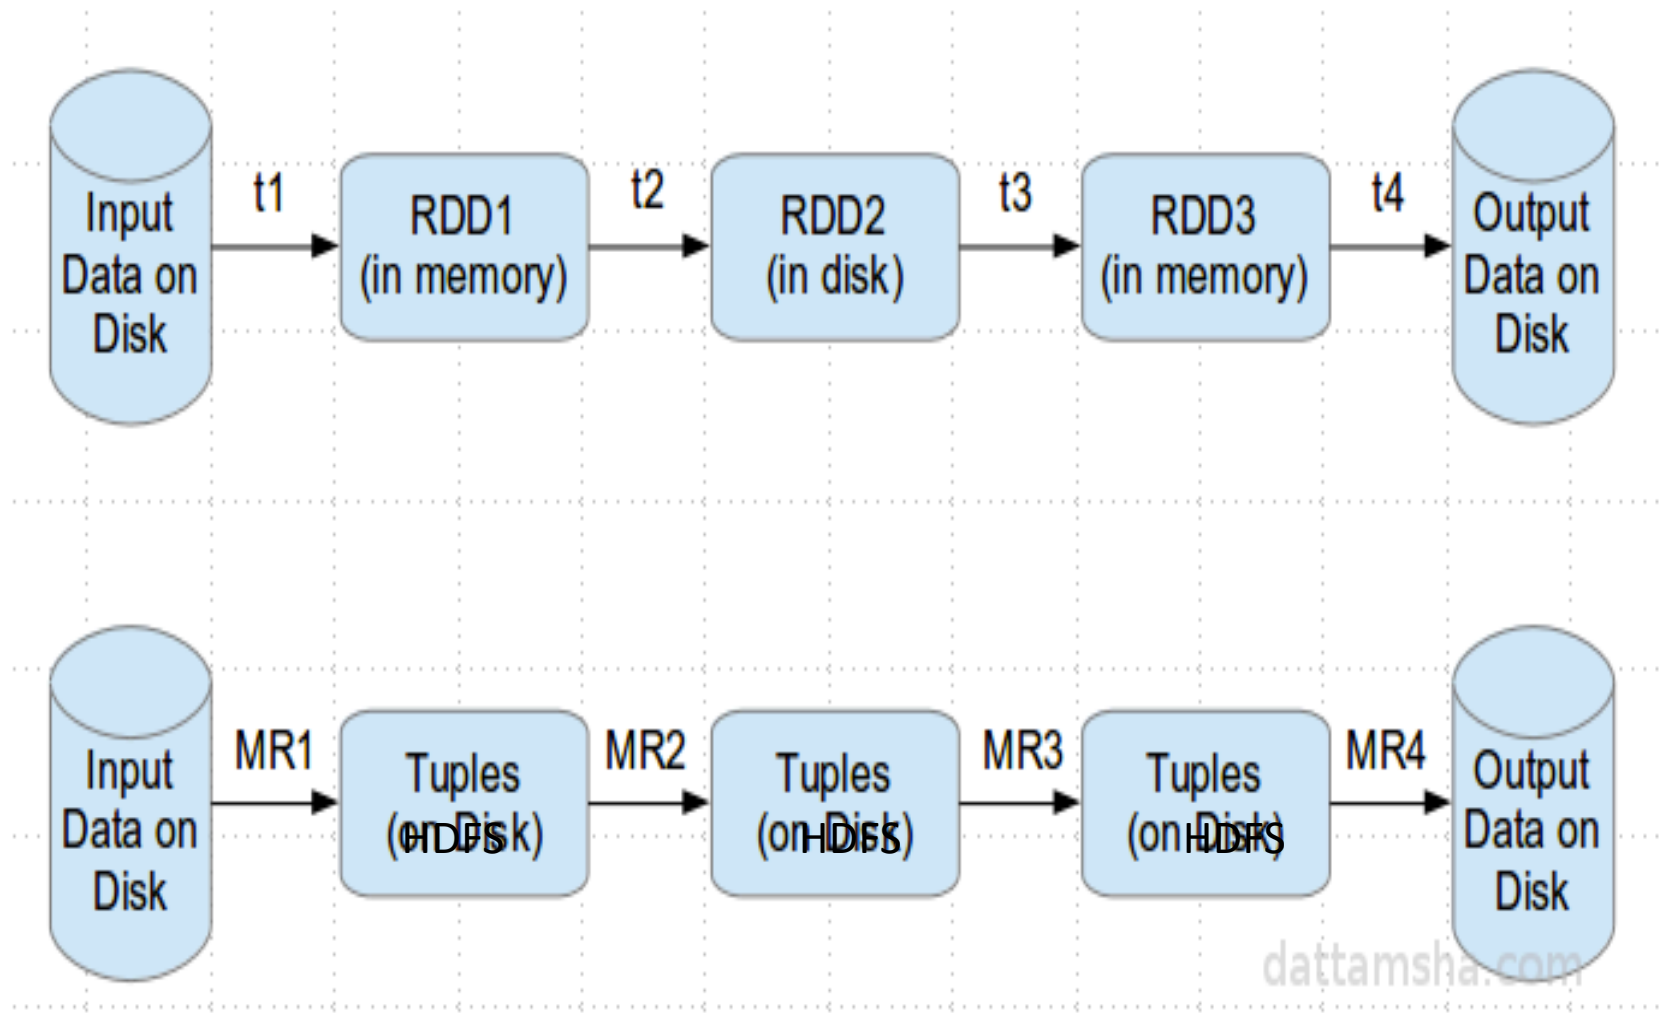


A unified analytics engine for large-scale data processing

Environments



Memory instead of disk



Spark and Map Reduce differences

	Apache Hadoop MR	Apache Spark
Storage	Disk only	In-memory or on disk
Operations	Map and Reduce	Many transformations and actions, including Map and Reduce
Execution model	Batch	Batch, iterative, streaming
Languages	Java	Scala, Java, Python and R

Apache Spark vs Apache Hadoop

	Hadoop World Record	Spark 100 TB *	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400	6592	6080
# Reducers	10,000	29,000	250,000
Rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min
Sort Benchmark Daytona Rules	Yes	Yes	No
Environment	dedicated data center	EC2 (i2.8xlarge)	EC2 (i2.8xlarge)

<https://databricks.com/blog/2014/10/10/spark-petabyte-sort.html>

Interactive shell

```
1. pyspark (java)
$ pyspark
Python 2.7.9 (default, Jan 7 2015, 11:49:12)
Type "copyright", "credits" or "license" for more information.

IPython 2.4.0 -- An enhanced Interactive Python.
?                -> Introduction and overview of IPython's features.
%quickref        -> Quick reference.
help             -> Python's own help system.
object?         -> Details about 'object', use 'object??' for extra details.
Welcome to

      _ _ _ _ _
     / _ _ _ \
    / _ _ _ \
   / _ _ _ \
  / _ _ _ \
 / _ _ _ \
/_ _ _ _ \
version 1.6.0




Using Python version 2.7.9 (default, Jan 7 2015 11:49:12)
SparkContext available as sc, HiveContext available as sqlContext.

In [1]:
```

(Scala, Python and Ronly)

Spark execution overview

Driver Program



Worker Machine



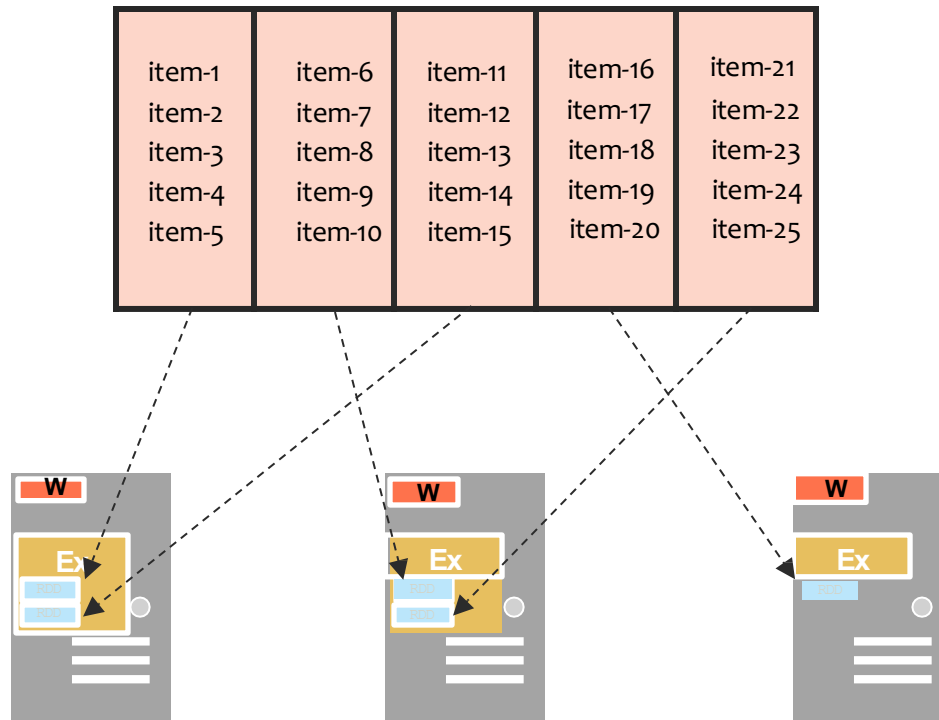
Worker Machine

Resilient Distributed Dataset (RDD)

- RDDs are *fault-tolerant, parallel data structures* that let users explicitly persist *intermediate results in memory*, control their partitioning to optimize data placement, and manipulate them using *a rich set of operators*.
 - RDDs automatically rebuilt on machine failure
- coarse-grained transformations vs. fine-grained updates
 - e.g., map, filter and join) that apply the same operation to many data items at once.

More partitions = more parallelism

RDD



RDD creation

- A base RDD can be created 2 ways:
 - Parallelize a collection
 - Read data from an external source (S3, C*, HDFS, etc)

Error, ts, msg1 Warn, ts, msg2 Error, ts, msg1	Info, ts, msg8 Warn, ts, msg2 Info, ts, msg8	Error, ts, msg3 Info, ts, msg5 Info, ts, msg5	Error, ts, msg4 Warn, ts, msg9 Error, ts, msg1
--	---	---	--

logLinesRDD

Parallelize



// Parallelize in Scala

```
val wordsRDD = sc.parallelize(List("fish", "cats", "dogs"))
```

- Take an existing in-memory collection and pass it to SparkContext's parallelize method

- Not generally used outside of prototyping and testing since it requires entire dataset in memory on one machine



Parallelize in Python

```
wordsRDD = sc.parallelize(["fish", "cats", "dogs"])
```



// Parallelize in Java

```
JavaRDD<String> wordsRDD = sc.parallelize(Arrays.asList("fish", "cats", "dogs"));
```

Read from text file



```
// Read a local txt file in Scala  
val linesRDD = sc.textFile("/path/to/README.md")
```

There are other methods to read data from HDFS, C*, S3, HBase, etc.



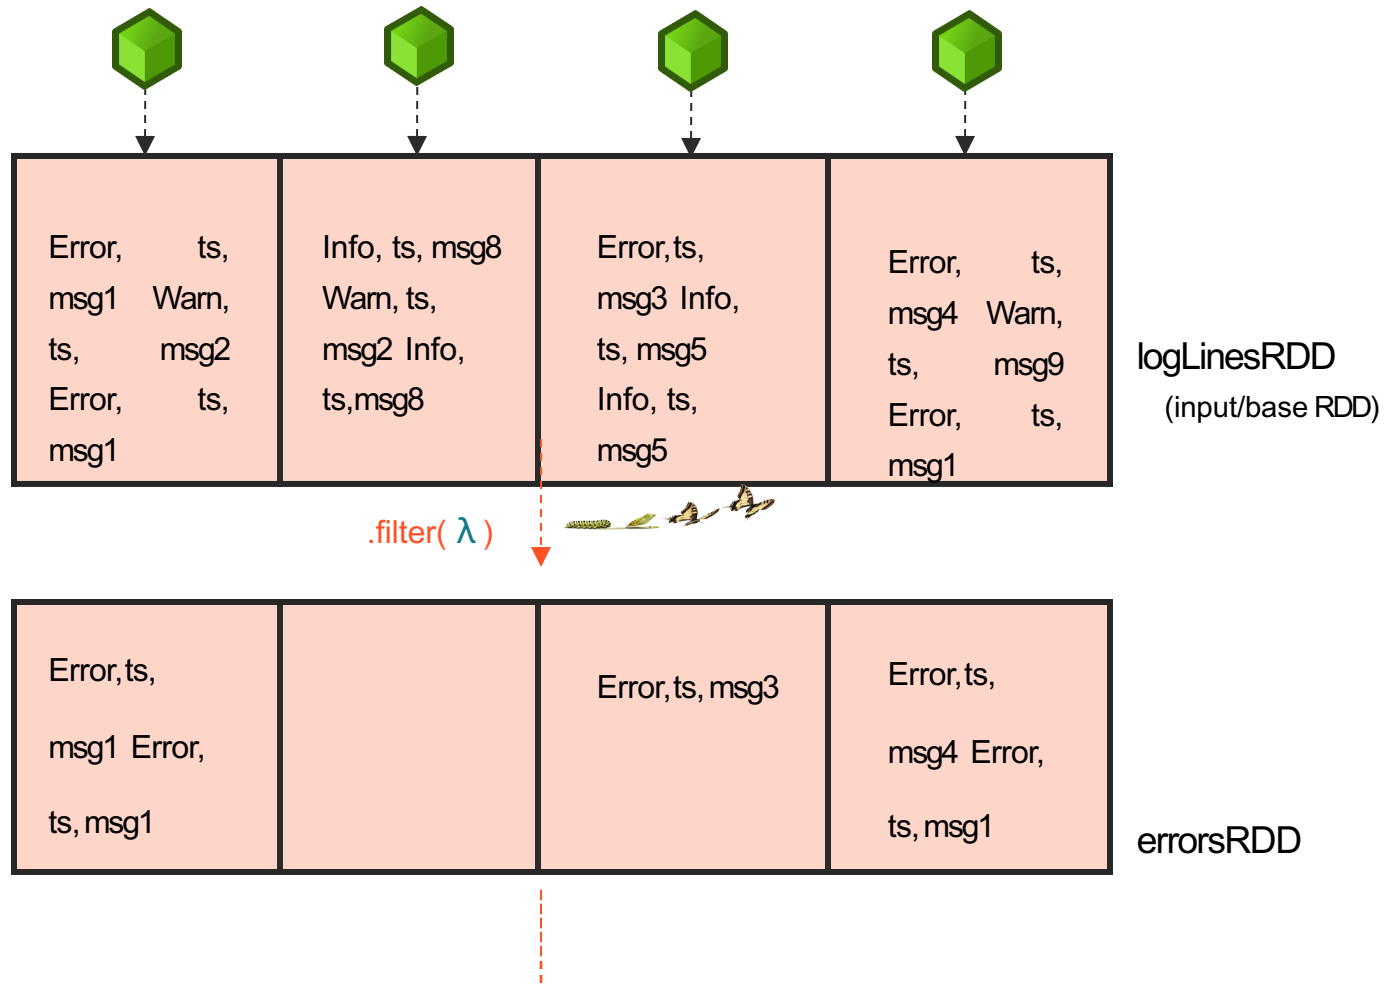
```
# Read a local txt file in Python  
linesRDD = sc.textFile("/path/to/README.md")
```

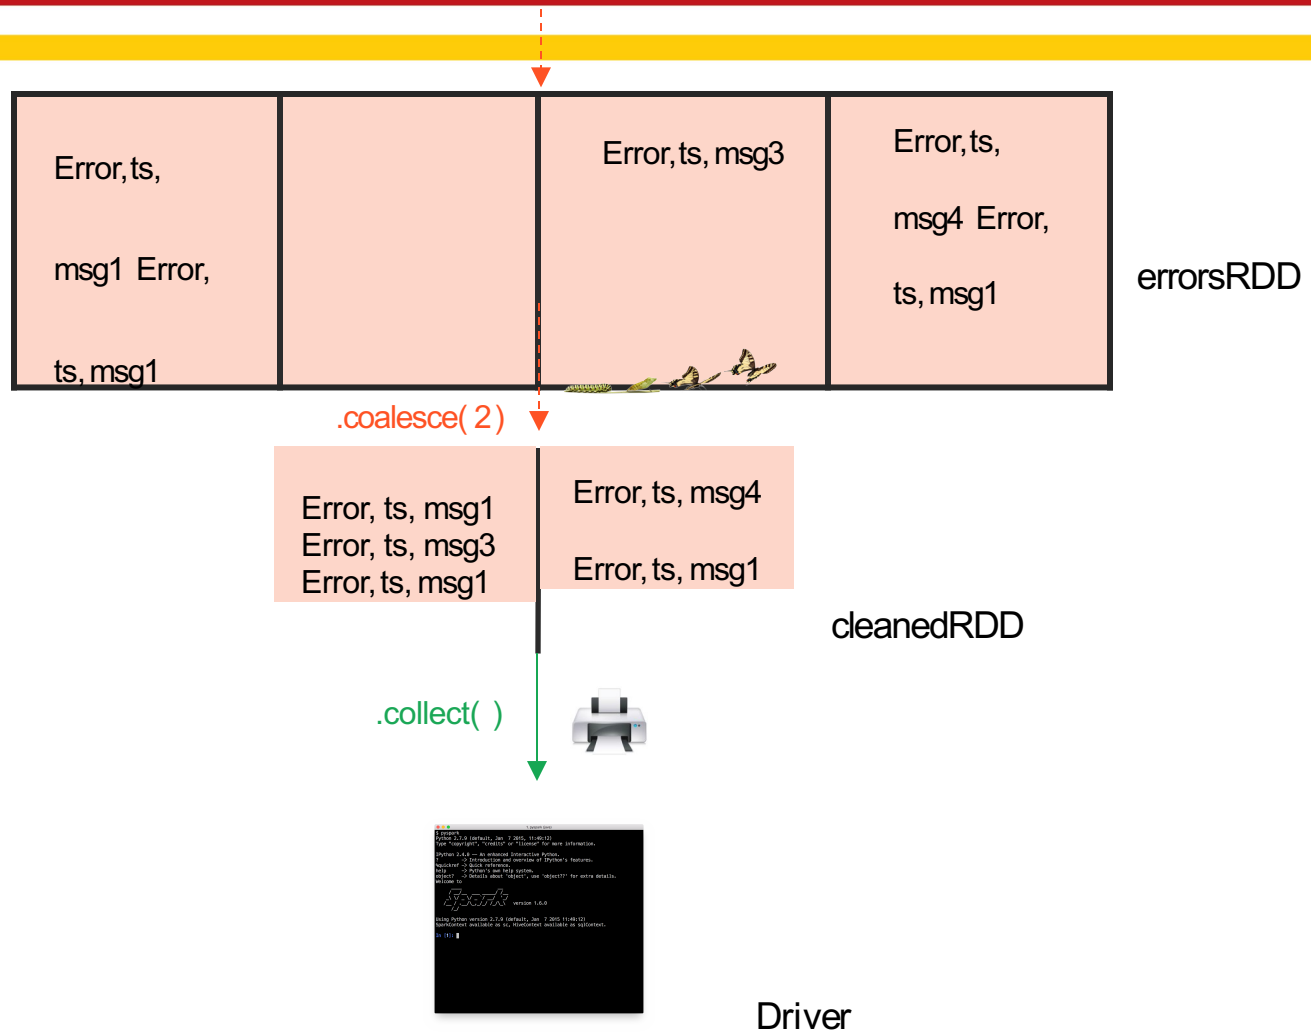


```
// Read a local txt file in Java  
JavaRDD<String> lines = sc.textFile("/path/to/README.md");
```


Operations on RDD

- Two types of operations: **transformations and actions**
- Transformations are lazy (not computed immediately)
- Transformations are executed when an action is run
- Persist (cache) distributed data in memory or disk








Execute DAG!

```

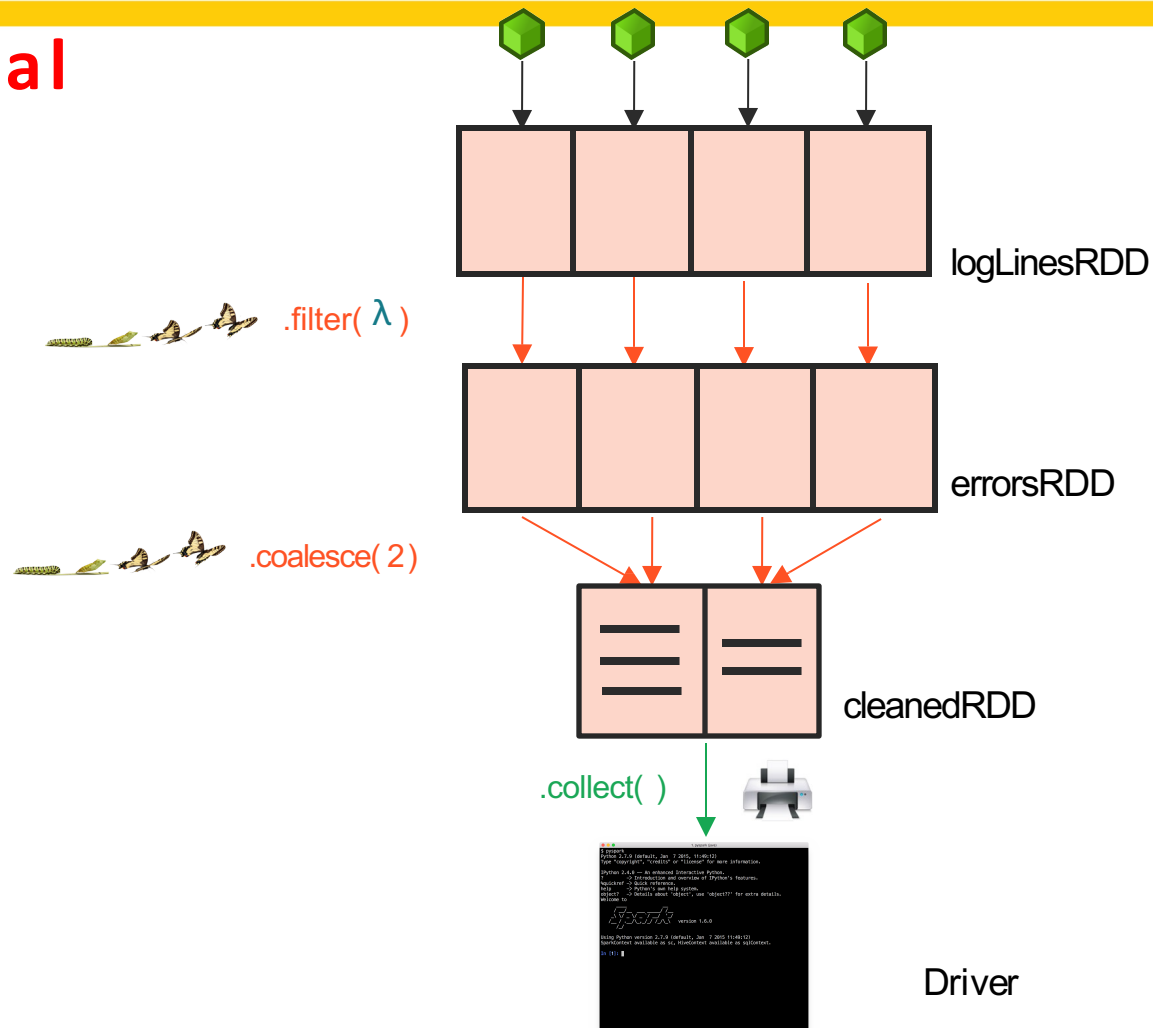
.collect( )

```




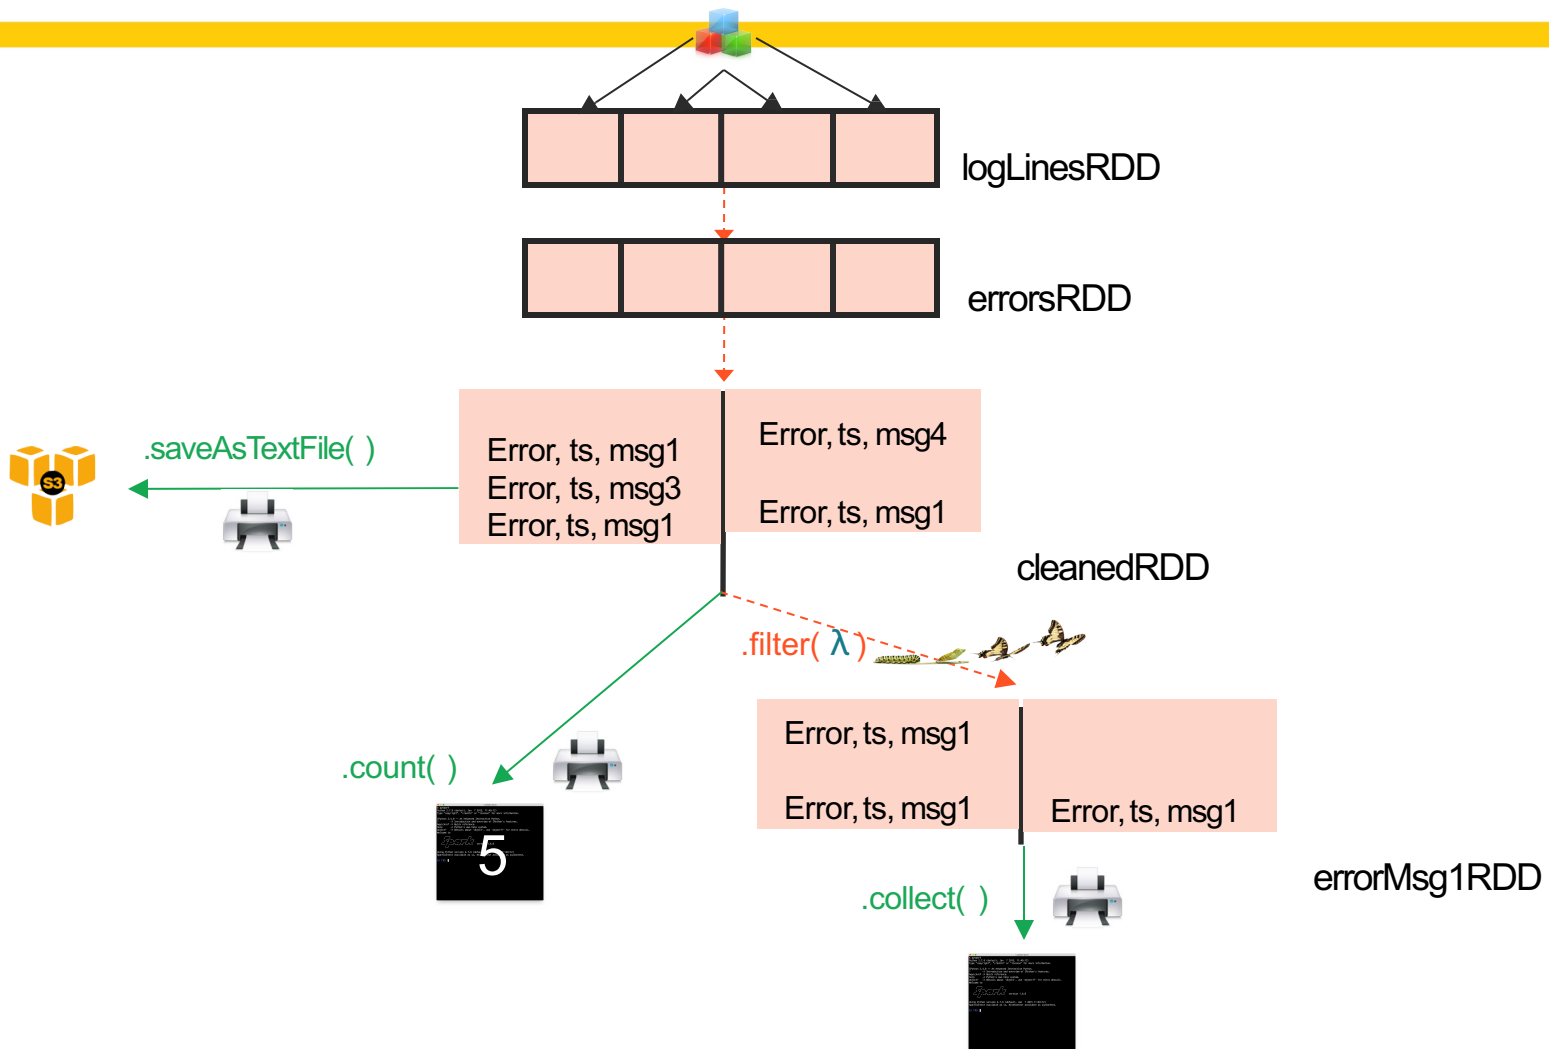
Driver

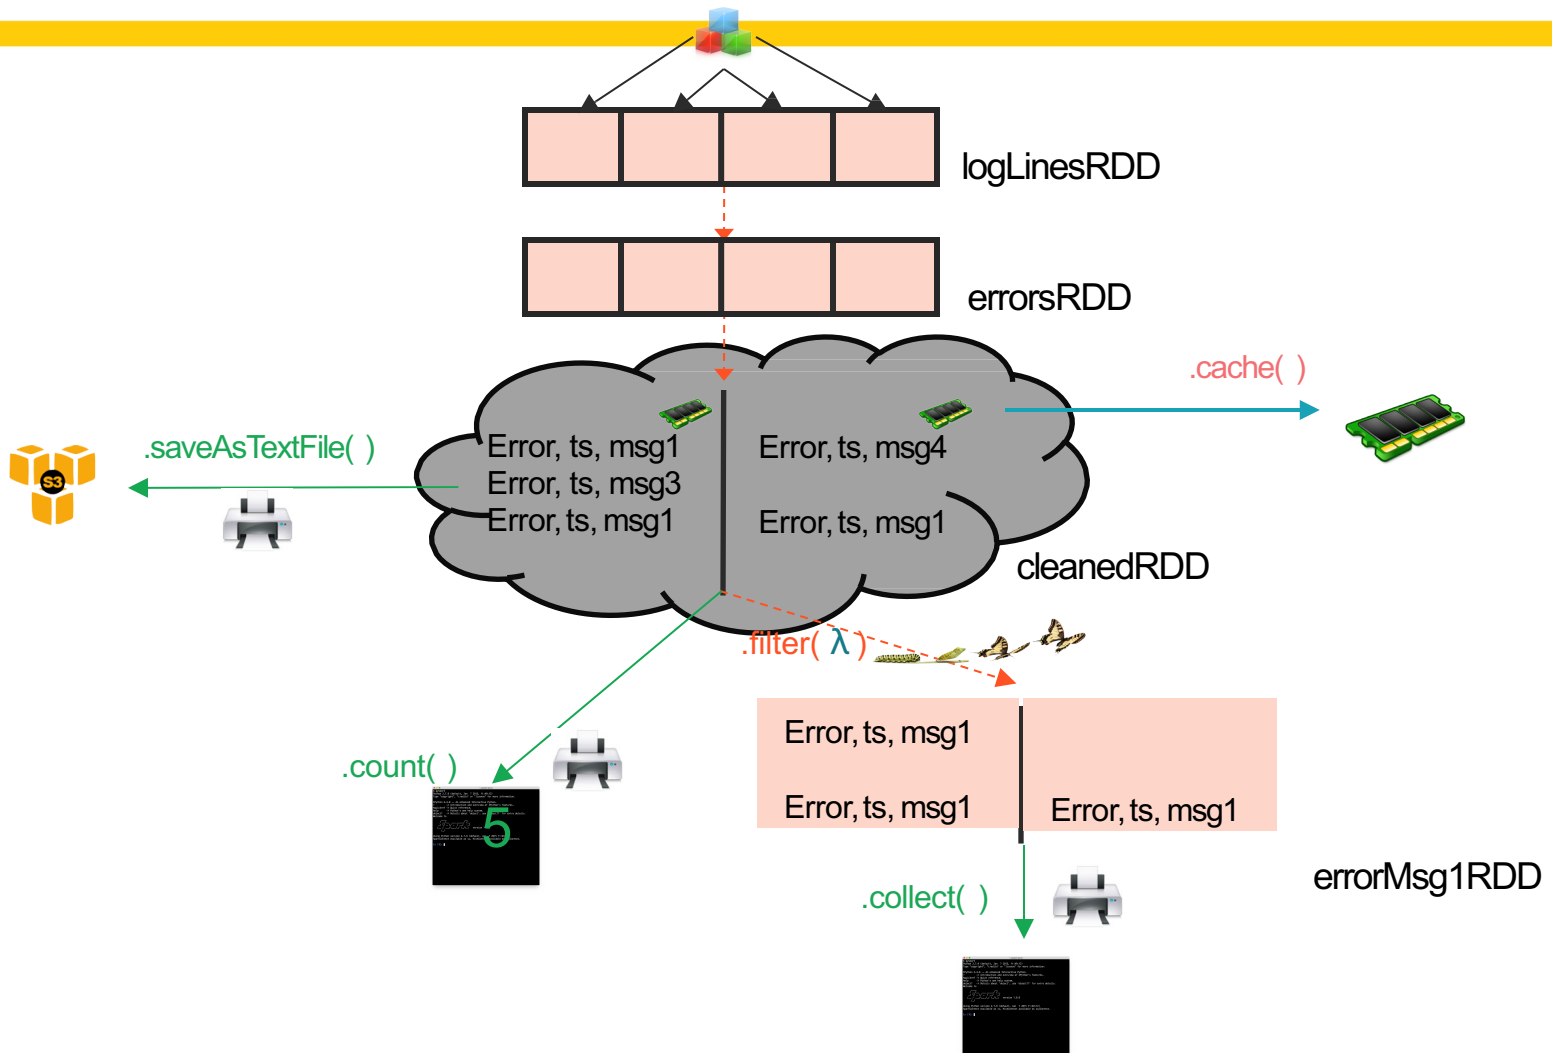
Logical



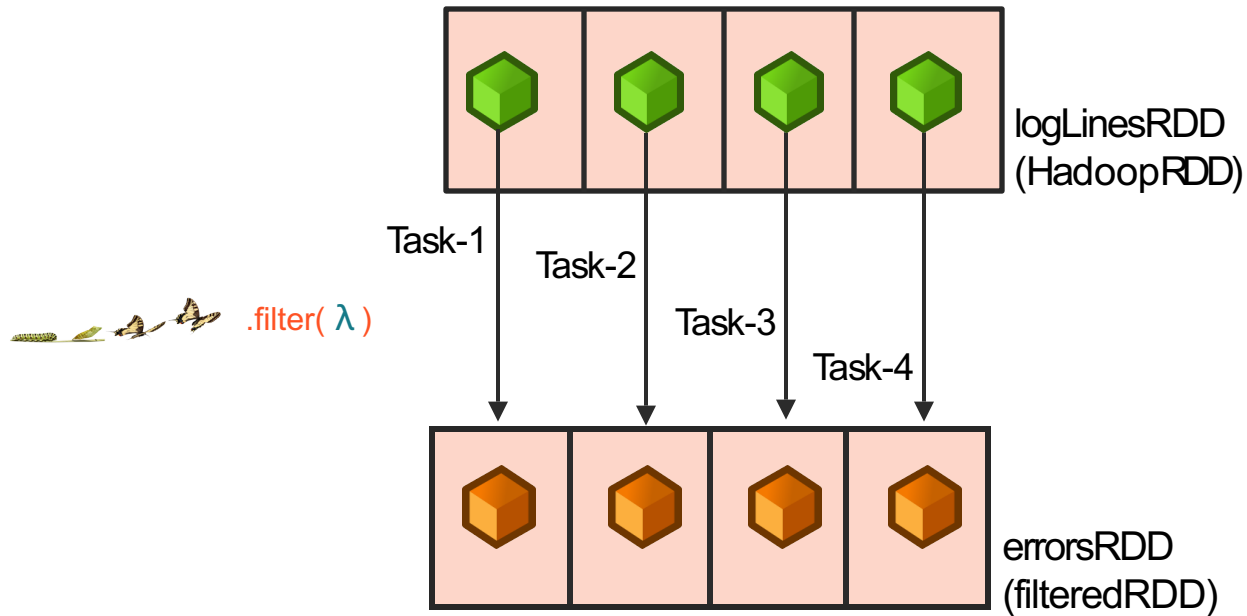
Physical



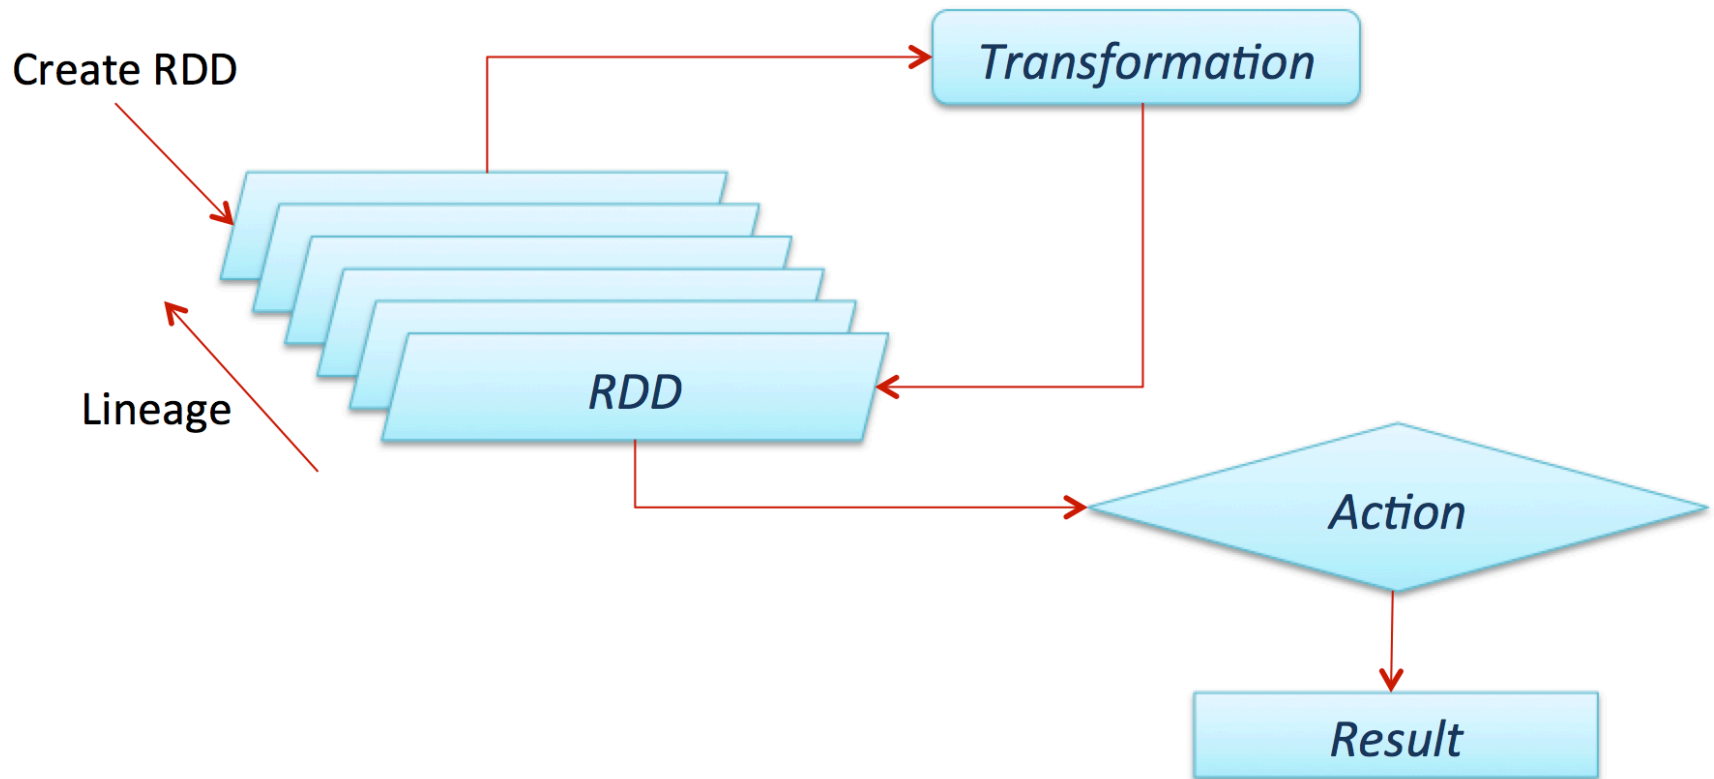




Partition >> Task >> Partition



Lineage mechanism



Notes on RDD

- Initial RDD on **disks** (HDFS, etc)
- Intermediate RDD on **RAM**
- Fault recovery based on **lineage**
- RDD operations is distributed

Spark program routine

- Create RDD from external data or create RDD from a collection in driver program
- Lazily transform them into new RDD
- `cache()` some RDD for reuse
- Perform actions to execute parallel computation and produce results

Transformations (lazy)

<code>map()</code>	<code>intersection()</code>	<code>cartesian()</code>
<code>flatMap()</code>	<code>distinct()</code>	<code>pipe()</code>
<code>filter()</code>	<code>groupByKey()</code>	<code>coalesce()</code>
<code>mapPartitions()</code>	<code>reduceByKey()</code>	<code>repartition()</code>
<code>mapPartitionsWithIndex()</code>	<code>sortByKey()</code>	<code>partitionBy()</code>
<code>sample()</code>	<code>join()</code>	<code>...</code>
<code>union()</code>	<code>cogroup()</code>	<code>...</code>

Actions

`reduce()`

`collect()`

`count()`

`first()`

`take()`

`takeSample()`

`saveToCassandra()`

`takeOrdered()`

`saveAsTextFile()`

`saveAsSequenceFile()`

`saveAsObjectFile()`

`countByKey()`

`foreach()`

...

Some Types of RDDs

- HadoopRDD
- FilteredRDD
- MappedRDD
- PairRDD
- ShuffledRDD
- UnionRDD
- PythonRDD
- DoubleRDD
- JdbcRDD
- JsonRDD
- VertexRDD
- EdgeRDD
- CassandraRDD
(DataStax)
- GeoRDD
(ESRI)
- EsSpark
(ElasticSearch)

Hands-on with RDDs



DataFrame

- A primary abstraction in Spark 2.0
 - Immutable once constructed
 - Track lineage information to efficiently re-compute lost data
 - Enable operations on collection of elements in parallel
- To construct DataFrame
 - By parallelizing existing Python collections (lists)
 - By transforming an existing Spark or pandas DataFrame
 - From files in HDFS or other storage system

Using DataFrame

- `>>> data = [('Alice', 1), ('Bob', 2), ('Bob', 2)]`
- `>>> df1 = sqlContext.createDataFrame(data, ['name', 'age'])`
- `[Row(name=u'Alice', age=1), Row(name=u'Bob', age=2),
Row(name=u'Bob', age=2)]`

Transformations

- Create new DataFrame from an existing one
- Use lazy evaluation
 - Nothing executes
 - Spark saves recipe for transformation source

Transformation	Description
<code>select(*cols)</code>	Selects columns from this DataFrame
<code>drop(col)</code>	Returns a new Dataframe that drops the specific column
<code>filter(func)</code>	Returns a new DataFrame formed by selecting those rows of the source on which <i>func</i> returns true
<code>where(func)</code>	Where is an alias for filter
<code>distinct()</code>	Returns a new DataFrame that contains the distinct rows of the source DataFrame
<code>sort(*cols, **kw)</code>	Returns a new DataFrame sorted by the specified columns and in the sort order specified by <i>kw</i>

Using Transformations

- `>>> data = [('Alice', 1), ('Bob', 2), ('Bob', 2)]`
- `>>> df1 = sqlContext.createDataFrame(data, ['name', 'age'])`
- `>>> df2 = df1.distinct()`
- `[Row(name=u'Alice', age=1), Row(name=u'Bob', age=2)]`
- `>>> df3 = df2.sort("age", ascending=False)`
- `[Row(name=u'Bob', age=2), Row(name=u'Alice', age=1)]`

Actions

- Cause Spark to execute recipe to transform source
- Mechanisms for getting results out of Spark

Action	Description
<code>show(<i>n</i>, <i>truncate</i>)</code>	Prints the first <i>n</i> rows of this DataFrame
<code>take(<i>n</i>)</code>	Returns the first <i>n</i> rows as a list of Row
<code>collect()</code>	Returns all the records as a list of Row (*)
<code>count()</code>	Returns the number of rows in this DataFrame
<code>describe(*<i>cols</i>)</code>	Exploratory Data Analysis function that computes statistics (count, mean, stddev, min, max) for numeric columns

Using Actions

- `>>> data = [('Alice', 1), ('Bob', 2)]`
- `>>> df = sqlContext.createDataFrame(data, ['name', 'age'])`
- `>>> df.collect()`
- `[Row(name=u'Alice', age=1), Row(name=u'Bob', age=2)]`
- `>>> df.count()`
- `2`
- `>>> df.show()`
- `+-----+-----+`
- `|name| age |`
- `+-----+-----+`
- `|Alice| 1|`
- `|Bob | 2|`
- `+-----+-----+`

Caching

- `>>> linesDF = sqlContext.read.text('...')`
- `>>> linesDF.cache()`
- `>>> commentsDF = linesDF.filter(isComment)`
- `>>> print linesDF.count(), commentsDF.count()`
- `>>> commentsDF.cache()`

Spark Programming Routine (Dataframe)

- Create DataFrames from external data or create DataFrame from a collection in driver program
- Lazily transform them into new DataFrames
- `cache()` some DataFrames for reuse
- Perform actions to execute parallel computation and produce results

Machine Learning Library (MLlib)

- 2 packages
 - spark.mllib
 - spark.ml
- ML algorithms
 - Common learning algorithms such as classification, regression, clustering and collaborative filtering
- Featurization
 - Feature extraction, transformations, dimensionality reduction and selection
- Utilities
 - Linear algebra, statistics, data handling, ...

ML: Transformer

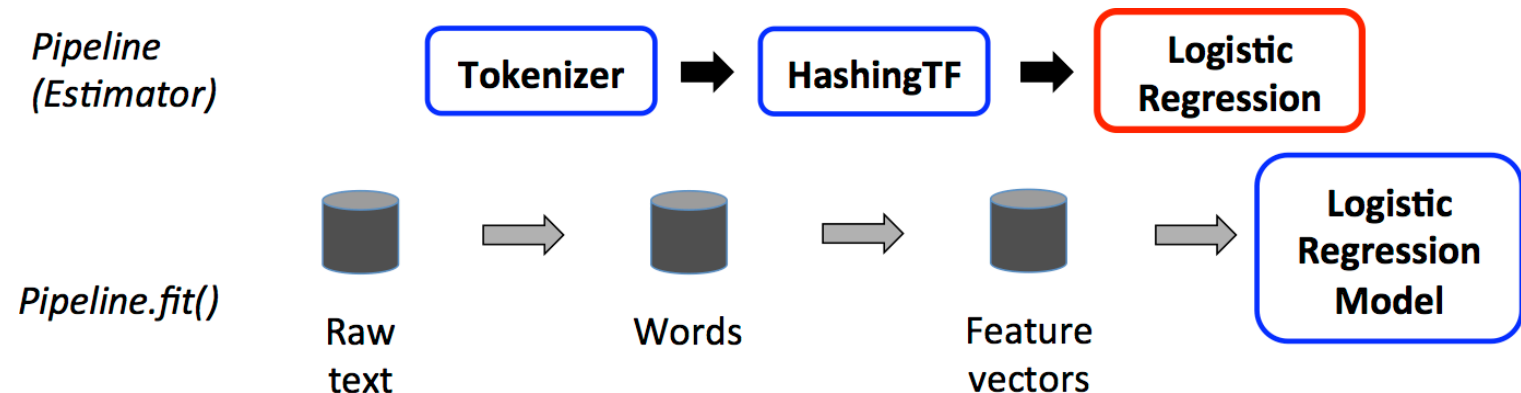
- A Transformer is a class which can transform a DataFrame into another DataFrame
- A Transformer implements transform()
- Examples
 - HashingTF
 - LogisticRegressionModel
 - Binarizer

ML: Estimator

- An Estimator is a class which can take a DataFrame and return a Transformer
- An Estimator implements fit()
- Examples
 - LogisticRegression
 - StandardScaler
 - Pipeline

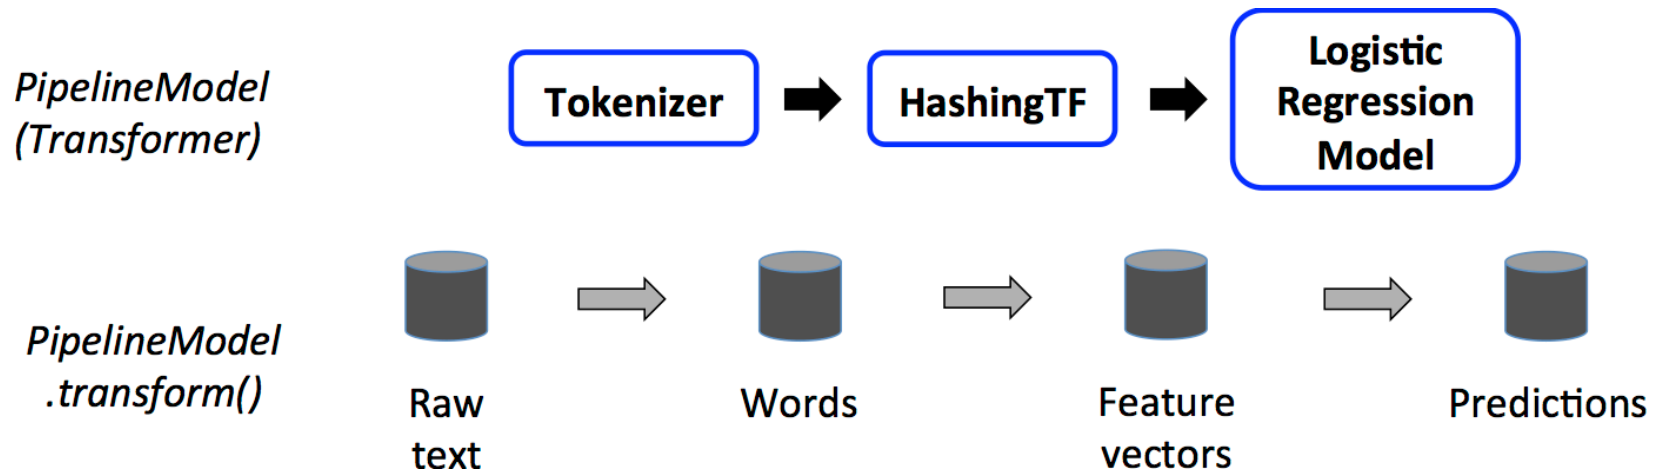
ML: Pipeline

- A Pipeline is an estimator that specified as a sequence of stages and each stage can be either estimators or transformers.




ML: PipelineModel

- After a `Pipeline.fit()` runs, it produces a `PipelineModel`. This `PipelineModel` is used at test time.



Architecture

- A master-worker type architecture
 - A driver or master node
 - Worker nodes



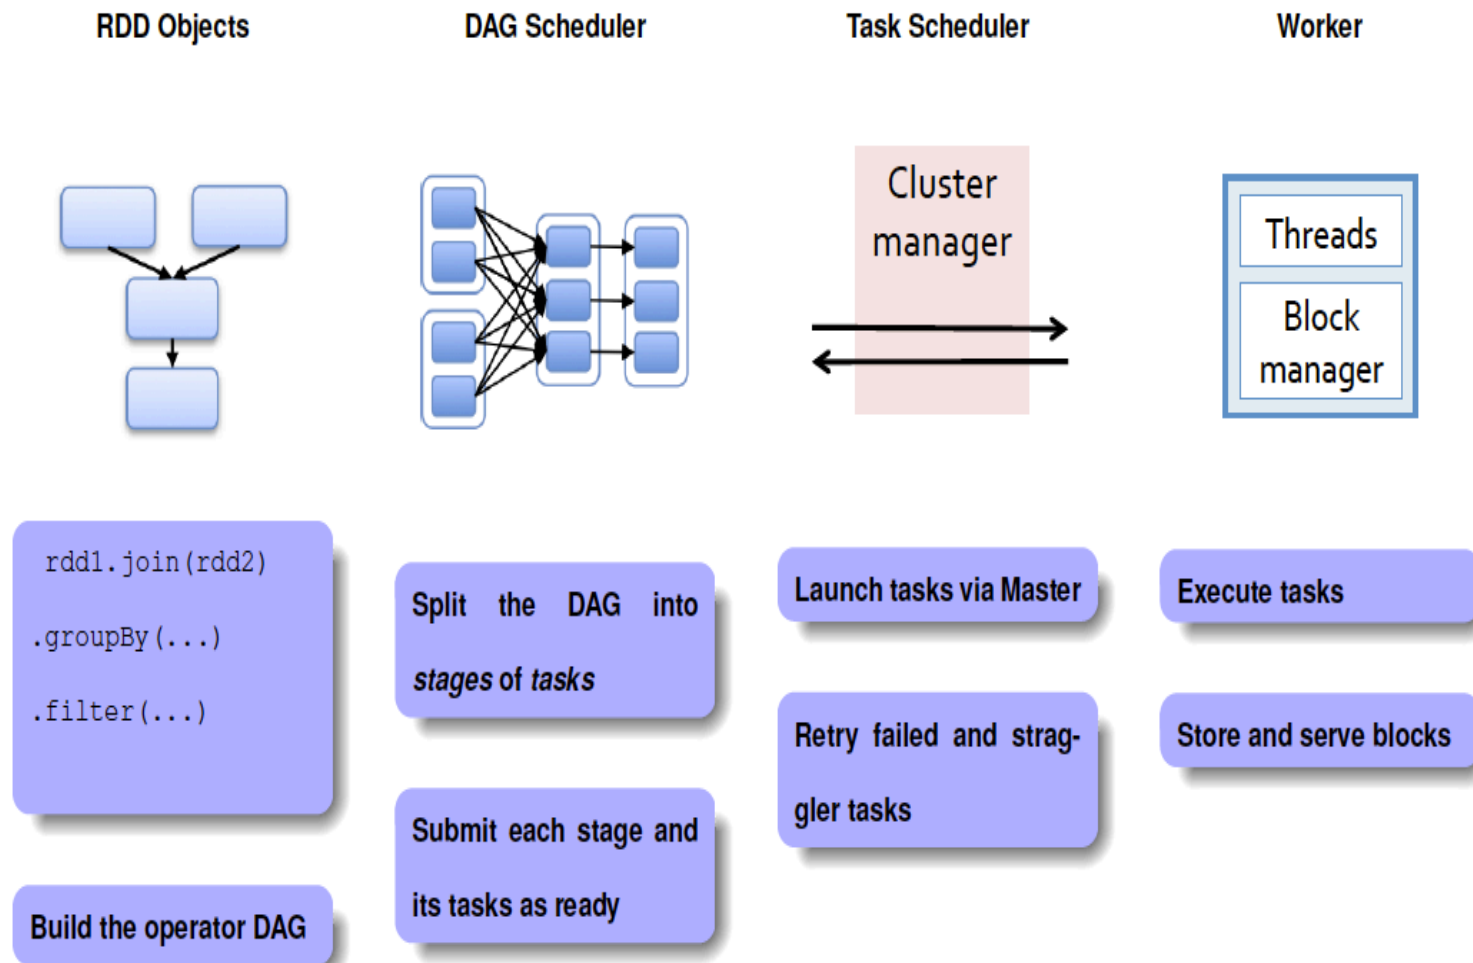
- The master send works to the workers and either instructs them to pull data from memory or from hard disk (or from another source like S3 or HDFS)

Architecture(2)

- A Spark program first creates a SparkContext object
 - SparkContext tells Spark how and where to access a cluster
 - The master parameter for a SparkContext determines which type and size of cluster to use

Master parameter	Description
local	Run Spark locally with one worker thread (no parallelism)
local[K]	Run Spark locally with K worker threads (ideal set to number of cores)
spark://HOST:PORT	Connect to a Spark standalone cluster
mesos://HOST:PORT	Connect to a Mesos cluster
yarn	Connect to a YARN cluster

Lifetime of a job in Spark




Narrow vs. Wide Dependencies

- Two types of transformations
 - Narrow transformation—doesn't require the data to be shuffled across the partitions. for example, Map, filter etc..
 - wide transformation—requires the data to be shuffled for example, reduceByKey etc..


narrow

each partition of the parent RDD is used by at most one partition of the child RDD




wide

multiple child RDD partitions may depend on a single parent RDD partition



Spark transformations & stages

- The narrow transformations will be grouped (pipe-lined) together into a single stage



Summary

- Hadoop: Scalable and economical data storage and processing
- Spark: a unified analytics engine for large-scale data processing



TRƯỜNG ĐẠI HỌC BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY OF SCIENCE AND TECHNOLOGY

Thank you for your attention!
Q&A

